

Generalized Linear Models


Mark Andrews

Psychology Department, Nottingham Trent University

✉ `mark.andrews@ntu.ac.uk`

The Poisson Distribution

- The Poisson distribution is a discrete probability distribution over the non-negative integers $0, 1, 2, \dots$



Shown here is a Poisson distribution with $\lambda = 3.5$.


The Poisson Distribution

- ▶ The Poisson distribution is used to model the probability of a given number of events occurring in a fixed interval of time, e.g. the number of emails you get per hour, the number of shark attacks on Bondi beach every summer, etc.
- ▶ It has a single parameter λ , known as the *rate*.
- ▶ If x is a Poisson random variable whose, its probability mass function is

$$P(x = k|\lambda) = \frac{e^{-\lambda}\lambda^k}{k!}.$$

The Poisson Distribution



- The mean of a Poisson distribution is equal to its rate parameter λ .
- Its variance is also equal to λ .



As λ increases, so too does the variance.

The Poisson Distribution

- The Poisson distribution can be seen as the limit of a Binomial distribution as $N \rightarrow \infty$, and $\lambda = pN$.
- Shown are (left) Binomial($N, p = \lambda/N$) where $N \approx 10^3$ and $\lambda = 7.5$, and (right) Poisson(λ).



Poisson Regression

- In any regression problem, our data are $(y_1, x_1), (y_2, x_2) \dots (y_n, x_n)$, where each y_i is modelled as a stochastic function of x_i .
- In Poisson regression, we assume that each y_i is a Poisson random variable rate λ_i and

$$\log(\lambda_i) = \beta_0 + \sum_{k=1}^K \beta_k x_{ki},$$

or equivalently

$$\lambda_i = e^{\beta_0 + \sum_{k=1}^K \beta_k x_{ki}}.$$

Poisson Regression

- ▶ As an example of Poisson regression, we can look at the number visits to a doctor in a fixed period as a function of predictors such as gender.
- ▶ Using a data-set of over 5000 people, we estimate (using mle) that

$$\log(\lambda_i) = 1.65 + 0.43 \times x_i$$

where $x_i = 1$ for a female, and $x_i = 0$ for a male.

Poisson Regression

- Using this example, we see that for a female

$$\lambda_{\text{Female}} = e^{1.65+0.43} = 8.004$$

and for males

$$\lambda_{\text{Male}} = e^{1.65} = 5.2$$

- In other words, the expected value for females is 8.2 and for males it is 5.2.

Coefficients

- ▶ In Poisson regression, coefficients can be understood as follows.
- ▶ In the previous example,

$$\begin{aligned}\lambda_{\text{Female}} &= e^{1.65+0.43}, \\ &= e^{1.65} e^{0.43}, \\ \lambda_{\text{Male}} &= e^{1.65}.\end{aligned}$$

- ▶ This means that the exponent of the gender coefficient, i.e. $e^{0.43}$, signifies the multiplicative increase to the average rate of doctor visits for women relative men.
- ▶ In other words, women visit doctors on average $e^{0.43} = 1.53$ times more than men.

Coefficients

- In an arbitrary example with a single continuous predictor variable,

$$\begin{aligned}\lambda &= e^{\alpha + \beta x_i}, \\ &= e^{\alpha} e^{\beta x_i},\end{aligned}$$

If we increase x_i by 1, we have

$$\begin{aligned}\lambda^+ &= e^{\alpha + \beta (x_i + 1)}, \\ &= e^{\alpha + \beta x_i + \beta}, \\ &= e^{\alpha} e^{\beta x_i} e^{\beta},\end{aligned}$$

- As $\lambda^+ = \lambda e^{\beta}$, we see that e^{β} is the multiplicative effect of an increase in one unit to the predictor variable.

Example

```
doc_df <- read_csv(here('data/DoctorAUS.csv')) %>%  
  mutate(gender = ifelse(sex == 1, 'female', 'male'))  
  
M <- glm(doctorco ~ gender,  
        data = doc_df,  
        family = poisson)  
  
doc_df_new <- tibble(gender = c('female', 'male'))  
  
doc_df_new %>%  
  add_predictions(M)  
  
doc_df_new %>%  
  add_predictions(M, type='response')
```

Model comparison

```
M_1 <- glm(doctorco ~ gender + insurance,  
           data = doc_df,  
           family = poisson)  
  
anova(M, M_1, test='Chisq')
```

Exposure and offset

- ▶ In some problems, the length of time during which events are measured varies across individuals.
- ▶ In the doctor visits example, we might have recordings of number of visits per year for some people and number of visits per 9 months, etc, for others.
- ▶ These situations are dealt with using an *exposure* term for each individual.

Exposure and offset

- ▶ When using an exposure term, we use the original count data as before, but treat

$$y_i \sim \text{Poisson}(\lambda_i/u_i),$$

where u_i is a term signifying the relative exposure time for person i .

- ▶ According to this,

$$\log(\lambda_i/u_i) = \alpha + \beta x_i,$$

$$\log(\lambda_i) = \alpha + \beta x_i + \log(u_i)$$

- ▶ In other words, $y_i \sim \text{Poisson}(\lambda_i/u_i)$ is equivalent to $y_i \sim \text{Poisson}(\lambda_i)$, where $\log(\lambda_i) = \alpha + \beta x_i + \log(u_i)$.

Exposure and offset

- ▶ For example, suppose we monitor people's drinking at social occasions. We find that three people drink 12, 7 and 3 drinks over the course of 7, 5 and 2 hours, respectively.
- ▶ If we are trying to predict drinking as a function of predictor variables, we ought to calibrate by the different time frames.
- ▶ Treating e.g. 12 as a draw from $\text{Poisson}(\lambda_i/7)$ where $\log(\lambda_i/7) = \alpha + \beta x_i$ is identical to treating 12 as a draw from $\text{Poisson}(\lambda_i)$ where $\log(\lambda_i) = \alpha + \beta x_i + \log(7)$.

Exposure and offset

- ▶ In general, exposure terms are treated as fixed offsets.
- ▶ If our data is $(y_1, x_1), (y_2, x_2) \dots (y_n, x_n)$ with exposures $u_1, u_2 \dots u_n$, then we treat

$$y_i \sim \text{Poisson}(\lambda_i),$$

where


$$\log(\lambda_i) = \log(u_i) + \beta_0 + \sum_{k=1}^K \beta_k x_{ki}.$$

Example

```
insur_df <- read_csv(here('data/Insurance.csv')) %>%  
  mutate(District = factor(District))  
  
M <- glm(Claims ~ District + Group + Age + offset(log(Holders)),  
  data = insur_df,  
  family = poisson)
```

The Poisson Distribution

- The mean of a Poisson distribution is equal to its rate parameter λ .
- Its variance is also equal to λ .



As λ increases, so too does the variance.

Means and variances in a Poisson distribution:

- ▶ In a Poisson distribution, the variance of a sample should be approximately the same as the mean of a sample.
- ▶ Example 1:

```
x <- rpois(25, lambda = 5)
c(mean(x), var(x), var(x)/mean(x))
#> [1] 5.4400000 5.4233333 0.9969363
```

- ▶ Example 2:


```
x <- rpois(25, lambda = 5)
c(mean(x), var(x), var(x)/mean(x))
#> [1] 5.4000000 6.416667 1.188272
```

Overdispersion

- ▶ If the variance of a sample is greater than would be expected according to a given theoretical model, then we say the data is *overdispersed*.
- ▶ In count data, if the variance of a sample is much greater than its mean, we say it is overdispersed.
- ▶ Using a Poisson distribution in this situation, this is an example of model mis-specification.
- ▶ It will also usually underestimate the standard errors in the Poisson model.

Overdispersion

- In the bioChemists data set, we have counts of the number of articles published by PhD students in the last three years (publications):



```
var(publications)/mean(publications)
#> [1] 2.191358
```

Overdispersion

- This leads standard errors to be *underestimated* if we use a Poisson model:

```
M <- glm(publications ~ 1, family=poisson)
summary(M)$coefficients
```

#>	Estimate	Std. Error	z value	Pr(> z)
#> (Intercept)	0.5264408	0.02540804	20.71945	2.312911e-95

Fixing overdispersion using a Quasi-poisson model

- ▶ A *quasi* Poisson model allows us to correct over-dispersion

```
M <- glm(publications ~ 1, family=quasipoisson)
summary(M)$coefficients
```

#>	Estimate	Std. Error	t value	Pr(> t)
#> (Intercept)	0.5264408	0.03761239	13.99647	1.791686e-40


- ▶ It does so by calculating an overdispersion parameter (roughly, the ratio of the variance to the mean) and multiplying the standard error by its square root.
- ▶ In this example, the overdispersion parameter is 2.1913892 and so its square root is 1.4803341.
- ▶ Alternatively, a *negative binomial regression* is an alternative to Poisson regression that can be used with overdispersed count data.

Negative binomial distribution

- ▶ A negative binomial distribution is a distribution over non-negative integers.
- ▶ To understand the negative binomial distribution, we start with the binomial distribution:
- ▶ If, for example, we have a coin whose probability of coming up heads is θ , then the number of Heads in a sequence of n flips will follow a binomial distribution.
- ▶ In this example, an outcome of Heads can be termed a *success*.


Negative binomial distribution

- Here is a binomial distribution where $n = 25$ and $\theta = 0.75$.




Negative binomial distribution

- ▶ A *negative* binomial distribution gives the probability distribution over the number of *successes* (e.g. Heads) before r *failures* (e.g. r Tails).
- ▶ For example, here we have the number of Heads (*successes*) that occur before we observe $r = 2$ Tails (*failures*), when the probability of Heads is $\theta = 0.75$:



Negative binomial distribution

- Here, we have the number of Heads (*successes*) that occur before we observe $r = 3$ Tails (*failures*), when the probability of Heads is $\theta = 0.5$:



Negative binomial distribution

- The probability mass function for the negative binomial distribution is:

$$P(x = k|r, \theta) = \binom{r+k-1}{k} \theta^r (1-\theta)^k$$

or more generally

$$P(x = k|r, \theta) = \frac{\Gamma(r+k)}{\Gamma(r)k!} \theta^r (1-\theta)^k,$$

where $\Gamma()$ is a Gamma function ($\Gamma(n) = (n-1)!$).

- In R, for any k , r , and θ , we can calculate $P(x = k|r, \theta)$ using `dnbinom`, e.g. $P(x = k = 2|r = 3, \theta = 0.75)$ is

```
dnbinom(2, 3, 0.75)  
#> [1] 0.1582031
```

Negative binomial distribution

- In the negative binomial distribution, the mean is

$$\mu = \frac{\theta}{1 - \theta} \times r,$$


and so

$$\theta = \frac{r}{r + \mu},$$

and we can generally parameterize the distribution by μ and r .

Why use negative binomial distribution?

- A negative binomial distribution is equivalent as weighted sum of Poissons.



- So it is appropriate to use when the data can be seen as arising from a mixture of Poisson distributions, each with different means.

Negative binomial regression

- In negative binomial regression, we have observed counts $y_1, y_2 \dots y_n$, and some predictor variables $x_1, x_2 \dots x_n$, and we assume that

$$y_i \sim \text{NegBinomial}(\mu_i, r),$$

where $\text{NegBinomial}(\mu_i, r)$ is a negative binomial with mean μ_i and a dispersion parameter r , and then

$$\log(\mu_i) = \beta_0 + \beta x_i.$$

Example

```
M <- glm.nb(publications ~ gender, data = biochemists_Df)
M1 <- glm.nb(publications ~ gender + married + I(children > 0),
```


Binomial logistic regression

- If y_i is the number of “successes” in n_i “trials”, we can model this as

$$y_i \sim \text{Binomial}(\theta_i, n_i),$$

$$\text{logit}(\theta_i) = \beta_0 + \sum_{k=1}^K \beta_k x_{ki}$$


- If $n_i = 1$ for all i , then this is exactly binary logistic regression.
- In general, it models the probability of something happening in a number of independent trials, and how the probability varies by the values of predictors.

Example

```
golf_df <- read_csv(here('data/golf_putts.csv')) %>%  
  mutate(failure = attempts - success,  
         p = success/attempts)  
  
M <- glm(cbind(success, failure) ~ distance,  
        family = binomial(link = 'logit'),  
        data = golf_df)
```


Poisson Distribution

A sample from a Poisson distribution with $\lambda = 5.5$.



Zero inflated Poisson Distribution

A sample from a zero inflated Poisson distribution with $\lambda = 5.5$, with probability of zero-component is 0.2.



Poisson regression to Zero-Inflated Poisson regression

- In Poisson regression (with a single predictor, for simplicity), we assume that each y_i is a Poisson random variable with rate λ_i that is a function of the predictor x_i .
- In Zero-Inflated Poisson regression, we assume that each y_i is distributed as a Zero-Inflated Poisson mixture model:

$$y_i \sim \begin{cases} \text{Poisson}(\lambda_i) & \text{if } z_i = 0, \\ 0, & \text{if } z_i = 1 \end{cases}$$

where rate λ_i and $P(z_i = 1)$ are functions of the predictor x_i .

Zero-Inflated Poisson regression

- Assuming data $\{(x_1, y_1), (x_2, y_2) \dots (x_n, y_n)\}$, Poisson regression models this data as:

$$y_i \sim \begin{cases} \text{Poisson}(\lambda_i) & \text{if } z_i = 0, \\ 0, & \text{if } z_i = 1 \end{cases},$$
$$z_i \sim \text{Bernoulli}(\theta_i),$$

where θ_i and λ_i are functions of x_i .

Zero-Inflated Poisson regression

- The θ_i and λ_i variables are the usual suspects, i.e.

$$\log(\lambda_i) = \alpha + \beta x_i,$$

and

$$\log\left(\frac{\theta_i}{1-\theta_i}\right) = a + b x_i.$$

- In other words, λ_i is modelled just as in ordinary Poisson regression and θ_i is modelled in logistic regression.

Examples

```
smoking_df <- read_csv(here('data/smoking.csv'))
M <- glm(cigs ~ educ, data = smoking_df)
M_zip <- zeroinfl(cigs ~ educ, data=smoking_df)

Df_new <- data.frame(educ = seq(20))
# Predicted average smoking rate
Df_new %>%
  add_predictions(M_zip, type='response')

# Predicted average smoking rate of "smokers"
Df_new %>%
  add_predictions(M_zip, type='count')

# Predicted probability of being a "smoker"
Df_new %>%
  add_predictions(M_zip, type='zero') %>%
  mutate(pred = 1-pred)
```